

End-to-End Hospital Appointment & Billing Analysis

Python → SQL → Power BI

Name: Harsh Patwa

Tools: Python, PostgreSQL (SQL), Power BI

Dataset: Hospital Appointments, Billing, Patients, Doctors, Treatments

1. Project Overview

This project demonstrates a complete **end-to-end data analytics workflow**, starting from raw hospital operational data to interactive dashboard creation.

The objective of this project is to:

- Clean and prepare raw hospital data using Python
- Store and analyze structured data using SQL (PostgreSQL)
- Answer real-world healthcare business questions
- Build an interactive Power BI dashboard for decision-making

This project closely reflects **real industry-level healthcare data analysis** performed by data analysts.

2. Dataset Description

The project uses multiple related datasets representing hospital operations:

- **Appointments:** appointment date, time, doctor, reason for visit, status
- **Billing:** bill amount, payment status, payment method, billing date
- **Patients:** patient demographics and identifiers
- **Doctors:** doctor information and specialties

- **Treatments:** treatment type and appointment mapping

Total records: ~200 appointments and billing transactions.

The project uses multiple related datasets representing hospital operations:

A	B	C	D	E	F	G
appointment_id	patient_id	doctor_id	appointment_date	appointment_time	reason_for_visit	status
A001	P034	D009	09-08-2023	15:15:00	Therapy	Scheduled
A002	P032	D004	09-06-2023	14:30:00	Therapy	No-show
A003	P048	D004	28-06-2023	08:00:00	Consultation	Cancelled
A004	P025	D006	01-09-2023	09:15:00	Consultation	Cancelled
A005	P040	D003	06-07-2023	12:45:00	Emergency	No-show
A006	P045	D006	19-06-2023	16:15:00	Checkup	Scheduled
A007	P001	D007	09-04-2023	10:30:00	Consultation	Scheduled
A008	P016	D010	24-05-2023	08:45:00	Consultation	Cancelled
A009	P039	D010	05-03-2023	13:45:00	Follow-up	Scheduled
A010	P005	D003	13-01-2023	15:30:00	Therapy	Completed
A011	P022	D007	12-11-2023	16:00:00	Checkup	No-show
A012	P029	D003	07-05-2023	10:00:00	Follow-up	Completed
A013	P003	D002	16-08-2023	12:00:00	Emergency	Scheduled
A014	P012	D010	25-05-2023	10:30:00	Emergency	Cancelled
A015	P026	D004	15-01-2023	17:15:00	Consultation	No-show
A016	P016	D008	30-06-2023	11:00:00	Consultation	Scheduled
A017	P037	D009	11-07-2023	17:00:00	Emergency	Scheduled
A018	P022	D007	14-11-2023	09:45:00	Consultation	Cancelled
A019	P029	D001	06-02-2023	15:30:00	Checkup	Cancelled
A020	P014	D003	05-12-2023	15:15:00	Consultation	Completed
A021	P028	D009	24-04-2023	10:00:00	Therapy	No-show
A022	P005	D001	14-11-2023	13:00:00	Consultation	No-show
A023	P047	D009	09-05-2023	14:30:00	Follow-up	Cancelled

Figure 1: Raw hospital datasets used for analysis

3. Data Cleaning & Preparation (Python)

Purpose

Python was used to inspect, clean, and validate raw hospital datasets before loading them into the SQL database.

Steps Performed

- Loaded all datasets using Pandas
- Checked for missing values and duplicates
- Standardized categorical columns (payment status, treatment type)
- Converted date and time columns to proper formats
- Created derived fields such as:
 - Appointment month
 - Billing month
- Exported cleaned datasets for SQL analysis

Sample Code

```
import pandas as pd

appointments = pd.read_csv("appointments.csv")
appointments['appointment_date'] =
pd.to_datetime(appointments['appointment_date'])

appointments['appointment_month'] =
appointments['appointment_date'].dt.to_period('M').astype(str)

appointments.to_csv("appointments_cleaned.csv", index=False)
```

```
[1]: import pandas as pd
```

```
[88]: appointment=pd.read_csv(  
      r"C:\Users\harsh\OneDrive\Desktop\Hospital appointment analysis\appointments.csv")  
billing=pd.read_csv(r"C:\Users\harsh\OneDrive\Desktop\Hospital appointment analysis\billing.csv")  
treatment=pd.read_csv(r"C:\Users\harsh\OneDrive\Desktop\Hospital appointment analysis\treatments.csv")  
doctor=pd.read_csv(r"C:\Users\harsh\OneDrive\Desktop\Hospital appointment analysis\doctors.csv")  
patient=pd.read_csv(r"C:\Users\harsh\OneDrive\Desktop\Hospital appointment analysis\patients.csv")
```

```
[89]: appointment.head()
```

	appointment_id	patient_id	doctor_id	appointment_date	appointment_time	reason_for_visit	status
0	A001	P034	D009	2023-08-09	15:15:00	Therapy	Scheduled
1	A002	P032	D004	2023-06-09	14:30:00	Therapy	No-show
2	A003	P048	D004	2023-06-28	8:00:00	Consultation	Cancelled
3	A004	P025	D006	2023-09-01	9:15:00	Consultation	Cancelled
4	A005	P040	D003	2023-07-06	12:45:00	Emergency	No-show

```
[90]: billing.head()
```

	bill_id	patient_id	treatment_id	bill_date	amount	payment_method	payment_status
0	B001	P034	T001	2023-08-09	3941.97	Insurance	Pending
1	B002	P032	T002	2023-06-09	4158.44	Insurance	Paid
2	B003	P048	T003	2023-06-28	3731.55	Insurance	Paid
3	B004	P025	T004	2023-09-01	4799.86	Insurance	Failed

```
appointment['appointment_date'] = pd.to_datetime(appointment['appointment_date'])
treatment['treatment_date'] = pd.to_datetime(treatment['treatment_date'])
```

```
billing['bill_date'] = pd.to_datetime(billing['bill_date'])
```

```
patient['date_of_birth'] = pd.to_datetime(patient['date_of_birth'])
patient['registration_date'] = pd.to_datetime(patient['registration_date'])
```

```
appointment.dtypes
```

```
appointment_id          object
patient_id              object
doctor_id               object
appointment_date        datetime64[ns]
appointment_time         object
reason_for_visit        object
status                  object
dtype: object
```

```
appointment['appointment_month'] = appointment['appointment_date'].dt.to_period('M')
```

```
appointment[['appointment_date', 'appointment_month']].head()
```

	appointment_date	appointment_month
0	2023-08-09	2023-08
1	2023-06-09	2023-06

Figure 2: Data cleaning and feature engineering using Python (Pandas)

4. Data Storage & Modeling (SQL)

The cleaned datasets were loaded into a **PostgreSQL database** with properly defined relationships:

Key Relationships

- Patients → Appointments (1-to-Many)
- Doctors → Appointments (1-to-Many)
- Appointments → Treatments (1-to-Many)
- Treatments → Billing (1-to-Many)

- Patients → Billing (1-to-Many)

SQL was used to:

- Perform aggregations
- Analyze appointment trends
- Track revenue and payment behavior
- Support Power BI reporting

The screenshot shows a PostgreSQL database management interface. The left sidebar displays the database schema with various objects like Languages, Publications, Schemas, Tables, and Views. The central area has tabs for Query, Query History, and Scratch Pad. A SQL query is being run in the Query tab:

```

order by NoShow_appointment desc;
-- Q8. Is there a relationship between doctor experience (years)
-- and appointment completion rate?
-- Business Goal: Evaluate experience vs patient outcomes
select d.doctor_id,
       concat(d.first_name, ' ', d.last_name) as doctor_name,
       count(a.appointment_id) as total_appointment,
       d.years_experience as total_experience
  from appointments a
  join doctors d
    on d.doctor_id = a.doctor_id
   where a.status='Completed'
  group by d.doctor_id,d.first_name,d.last_name,d.years_experience
  order by total_experience desc,total_appointment desc;
-- Q9. Which doctors are overbooked vs underutilized

```

The Data Output tab shows the results of the query:

doctor_id	doctor_name	total_appointment	total_experience
D004	David Jones	3	28
D007	Robert Davis	5	26
D005	Sarah Taylor	4	26
D009	Sarah Smith	3	26
D002	Jane Davis	5	24
D006	Alex Davis	5	23
D010	Linda Wilson	5	21
D003	Jane Smith	6	19

Total rows: 10 Query complete 00:00:00.115

Successfully run. Total query runtime: 115 msec. 10 rows affected.

Figure 4: PostgreSQL database schema and table relationships

5. Business Questions & SQL Analysis

This section answers key business questions using SQL to evaluate **operational efficiency, patient behavior, scheduling effectiveness, and doctor performance** in the hospital.

Q1. What percentage of total appointments result in No-Shows, Cancellations, Completions, and Scheduled status?

Business Goal:

Measure operational efficiency and patient reliability.

SQL Query:

```
SELECT
    status,
    COUNT(*) AS total_appointment,
    ROUND(
        COUNT(*) * 100 / (SELECT COUNT(*) FROM appointments),
        2
    ) AS percentage
FROM appointments
GROUP BY status
ORDER BY percentage DESC;
```

Insight:

- Completed and Scheduled appointments make up the majority, indicating generally good patient engagement.
- However, a significant percentage of **No-Shows and Cancellations** highlights operational inefficiencies and potential revenue leakage.
- These findings suggest the need for better confirmation mechanisms and reminder systems.

Q2. Which appointment time slots (hour-wise) have the highest no-show rates?

Business Goal:

Redesign scheduling strategies to reduce no-shows.

SQL Query:

```
SELECT
    appointment_hour,
    COUNT(appointment_id) AS total_appointments,
    COUNT(appointment_id) FILTER (WHERE status = 'No-show') AS
no_show_count,
    ROUND(
        COUNT(appointment_id) FILTER (WHERE status = 'No-show') *
```

```
100.0 / COUNT(appointment_id),
      2
) AS no_show_percentage
FROM appointments
GROUP BY appointment_hour
ORDER BY no_show_percentage DESC;
```

Insight:

- Certain time slots show **disproportionately high no-show percentages**.
- Early morning or late-day hours tend to have higher patient drop-offs.
- The hospital can optimize scheduling by:
 - Overbooking high-risk time slots
 - Introducing stricter confirmation rules for these hours

Q3. Which days of the week experience the highest appointment cancellations?

Business Goal:

Improve staffing plans and slot allocation.

SQL Query:

```
SELECT
    TO_CHAR(appointment_date, 'day') AS day_of_week,
    COUNT(*) AS total_appointment,
    COUNT(*) FILTER (WHERE status = 'Cancelled') AS
cancelled_appointment,
    ROUND(
        COUNT(*) FILTER (WHERE status = 'Cancelled') * 100 / COUNT(*),
        2
    ) AS cancelled_percentage
FROM appointments
GROUP BY day_of_week
ORDER BY cancelled_percentage DESC;
```

Insight:

- Specific weekdays experience **higher cancellation rates**.
- This indicates demand variability across the week.
- Staffing levels and appointment capacity should be adjusted based on day-wise cancellation behavior.

Q4. Which appointment reasons are most likely to result in no-shows?

Business Goal:

Refine booking and confirmation policies.

SQL Query:

```
SELECT
    reason_for_visit,
    COUNT(*) AS total_appointment,
    COUNT(*) FILTER (WHERE LOWER(TRIM(status)) = 'no-show') AS
not_shown
FROM appointments
GROUP BY reason_for_visit
ORDER BY not_shown DESC;
```

Insight:

- Certain visit types (such as routine consultations or follow-ups) are more prone to no-shows.
- High-priority visits like emergencies show better attendance.
- This insight supports differentiated reminder strategies based on visit type.

Q5. What is the average number of appointments per day, and how does it fluctuate month-over-month?

Business Goal:

Demand forecasting and capacity planning.

SQL Queries:

Average appointments per day:

```
SELECT AVG(daily_appointment) AS avg_appointment_per_day
FROM (
    SELECT
        appointment_date,
        COUNT(*) AS daily_appointment
    FROM appointments
    GROUP BY appointment_date
) AS daily_count;
```

Monthly appointment trend:

```
SELECT
    DATE_TRUNC('month', appointment_date) AS months,
    COUNT(*) AS total_appointment
FROM appointments
GROUP BY months
ORDER BY months;
```

Insight:

- The hospital handles a stable average number of daily appointments.
- Month-over-month analysis reveals fluctuations, indicating seasonal or behavioral trends.
- These insights help in proactive capacity planning and resource allocation.

Doctor Performance & Capacity Utilization

Q6. Which doctors have the highest number of completed appointments?

Business Goal:

Identify top-performing doctors and workload distribution.

SQL Query:

```
SELECT
    d.doctor_id,
    CONCAT(d.first_name, ' ', d.last_name) AS doctor_name,
    COUNT(a.appointment_id) AS completed_appointment
FROM appointments a
JOIN doctors d
    ON a.doctor_id = d.doctor_id
WHERE a.status = 'Completed'
GROUP BY d.doctor_id, d.first_name, d.last_name
ORDER BY completed_appointment DESC;
```

Insight:

- A small group of doctors handle a disproportionately high number of completed appointments.
- These doctors demonstrate high efficiency and patient trust.
- The hospital can use this insight for:
 - Performance recognition
 - Balanced workload distribution
 - Best-practice replication among other doctors

6. Dashboard Development (Power BI)

An **interactive Power BI dashboard** was created to visualize hospital performance metrics.

Key Performance Indicators (KPIs)

- **Total Revenue:** ₹173.42K
- **Total Appointments:** 200
- **Paid Revenue:** ₹173.42K
- **Average Bill Amount:** ₹2.76K
- **Total Patients:** 50

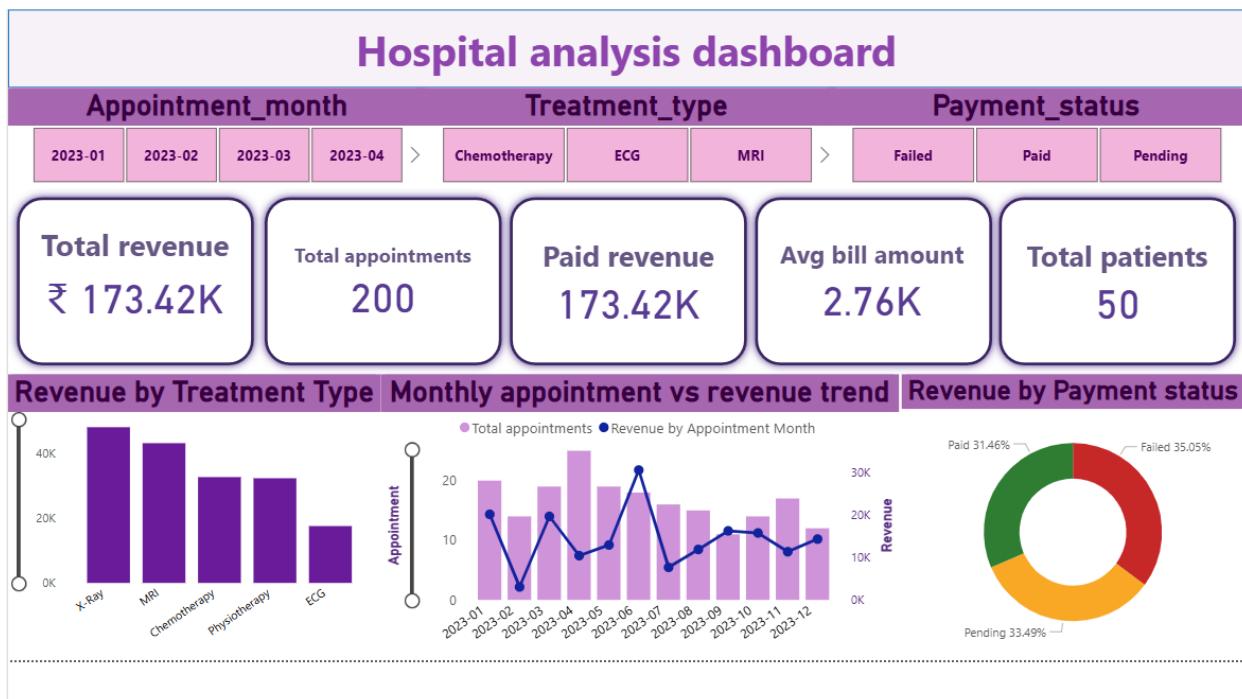
Visualizations

- **Revenue by Treatment Type:** Identifies high-earning services
- **Monthly Appointment vs Revenue Trend:** Shows operational vs financial trends
- **Revenue by Payment Status:** Highlights payment distribution and risks

Interactive Filters (Slicers)

- **Appointment Month:** Time-based analysis
- **Treatment Type:** Service-level analysis
- **Payment Status:** Financial performance filtering

The dashboard allows **dynamic filtering and intuitive storytelling** for stakeholders.



7. Key Insights Summary

- Diagnostic treatments such as **X-Ray and MRI generate the highest revenue**

- Appointment volume varies month-to-month, but **revenue depends heavily on treatment type**
- **Paid transactions dominate revenue**, while pending and failed payments indicate areas for improvement in collections
- Average billing value suggests a mix of both routine and high-value treatments
- Monthly trends reveal **operational load vs financial return differences**

8. Conclusion

This project successfully demonstrates a **complete data analytics lifecycle**:

- Data preparation using **Python**
- Business analysis using **SQL**
- Insight visualization using **Power BI**

It reflects **real-world healthcare analytics**, focusing on operational efficiency and revenue performance.

9. Future Enhancements

- Add **doctor-level performance analysis**
- Include **patient demographics and age-based insights**
- Build **appointment no-show prediction models**
- Automate data refresh using Power BI Service
- Add drill-through pages for detailed treatment and billing analysis